



# Introduction to High Performance Computing

MPI (1)



# Motivation

OpenMP is only available on systems with shared memory.  
The parallelism is limited to the number of cores in ONE system.

Parallelism among different nodes requires explicit communication between tasks.  
Messages have to be exchanged.

## Message Passing Interface

- standardized parallelization library since 1994
- very performant (almost Hardware latency/bandwidth)
- many different implementations (different focus, but conforming to standard)
- still under development to incorporate multicore architecture (hybrid with OpenMP)
- development needed to incorporate manycore architectures

# MPI-versions

## **MPI-versions**

### **Intel-MPI**

module add intel/latest

mpicc for compilation (calls icc (option `-cc=icc`), binds appropriate libraries, header, ...)

**`bsub mpiexec.hydra ./exe`**

to execute in Batch-mode – additional options automatically added by LSF

more Information: [Software](#) page on Elwetritsch

# MPI-versions

## Platform-MPI

`module add platform-mpi/latest`

`mpicc` for compilation (calls gcc, binds appropriate libraries, header, ...)

`bsub -a platformmpi mpirun -lsb_mcpu_hosts ./exe`  
to execute – additional options automatically added by LSF

more Information: [Software](#) page on Elwetritsch

# MPI-versions

**OpenMPI** (<http://www.open-mpi.org>)

```
module add intel-2015
```

```
module add openmpi/1.8.6-intel-2015
```

mpicc for compilation (calls icc (option `-cc=icc`), binds appropriate libraries, header, ...)

```
bsub -a openmpi mpirun ./exe
```

to execute – additional options automatically added by LSF

more Information: [Software](#) page on Elwetritsch

# Interfaces

Simple script to compile and/or link MPI programs.

Usage: mpigcc [options] <files>

-----  
The following options are supported:

-cc=<name> specify a C compiler name: i.e. -cc=gcc  
-echo print the scripts during their execution  
-show show command lines without real calling  
-config=<name> specify a configuration file: i.e. -config=gcc for mpicc-gcc.conf file  
-v print version info of mpigcc and its native compiler  
-profile=<name> specify a profile configuration file (an MPI profiling library): i.e. -profile=myprofile for the myprofile.cfg file.  
As a special case, lib<name>.so or lib<name>.a may be used if the library is found  
-check\_mpi link against the Intel(R) Trace Collector (-profile=vtmc).  
-static\_mpi link the Intel(R) MPI Library statically  
-mt\_mpi link the thread safe version of the Intel(R) MPI Library  
-ilp64 link the ILP64 support of the Intel(R) MPI Library  
-t or -trace link against the Intel(R) Trace Collector  
-dynamic\_log link against the Intel(R) Trace Collector dynamically  
-static use static linkage method  
-nostrip turn off the debug information stripping during static linking  
-O enable optimization  
-link\_mpi=<name> link against the specified version of the Intel(R) MPI Library  
All other options will be passed to the compiler without changing.

-----  
The following environment variables are used:

I\_MPI\_ROOT the Intel(R) MPI Library installation directory path  
I\_MPI\_CC or MPICH\_CC the path/name of the underlying compiler to be used  
I\_MPI\_CC\_PROFILE or MPICC\_PROFILE the name of profile file (without extension)  
I\_MPI\_COMPILER\_CONFIG\_DIR the folder which contains configuration files \*.conf  
I\_MPI\_TRACE\_PROFILE specify a default profile for the -trace option  
I\_MPI\_CHECK\_PROFILE specify a default profile for the -check\_mpi option  
I\_MPI\_CHECK\_COMPILER enable compiler setup checks  
I\_MPI\_LINK specify the version of the Intel(R) MPI Library  
I\_MPI\_DEBUG\_INFO\_STRIP turn on/off the debug information stripping during static linking  
-----



# Interfaces

```
schuele@head3 [~] mpirun -help
```

```
Usage: ./mpiexec [global opts] [exec1 local opts] : [exec2 local opts] : ...
```

```
Global options (passed to all executables):
```

green: usually not  
used

```
Global environment options:
```

```
Other global options:
```

```
-f {name} | -hostfile {name}    file containing the host names  
-hosts {host list}             comma separated host list  
-configfile {name}            config file containing MPMD launch options  
-machine {name} | -machinefile {name}  file mapping procs to machines
```

```
Local options (passed to individual executables):
```

```
Local environment options:
```

```
Other local options:
```

automatically set by batch  
system

```
Hydra specific options (treated as global):
```

```
Bootstrap options:
```

```
Resource management kernel options:
```

```
Processor topology options:
```

```
-binding                process-to-core binding mode
```

```
Checkpoint/Restart options:
```

```
Demux engine options:
```

```
Debugger support options:
```

```
-tv                    run processes under TotalView
```

```
Other Hydra options:
```

```
-perhost <n>           place consecutive <n> processes on each host  
-ppn <n>              stand for "process per node"; an alias to -perhost <n>
```

# Example

```
#include "mpi.h"
/* Simple Example for two Tasks */

main(int argc, char **argv)
{
    char text[20];
    int  myrank, size, sender=0, adressat=1, tag=99;
    MPI_Status status;

    MPI_Init(&argc, &argv);                /* Initialization */
    /* get unique ID*/
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    /* get number of tasks involved */
    MPI_Comm_size(MPI_COMM_WORLD, &size);
```



# Example

```
if(size != 2) {
    printf("Example for just 2 Tasks\n");
    MPI_Finalize();      exit(1);  }
if(myrank == 0) { /* different tasks do different things*/
    strcpy(text,"Hallo there");
    MPI_Send(text,strlen(text),MPI_CHAR,adressat,
              tag,MPI_COMM_WORLD);          /* Send */
} else {
    MPI_Recv(text,20,MPI_CHAR,sender,
              tag,MPI_COMM_WORLD,&status);/* Receive */
    printf("Task %d received:%s:\n",myrank,text);
}

MPI_Finalize();          /* End*/
}
```



# Example


- All functions use MPI\_X (capital), rest small
- Tasks are available after MPI\_Init
- Communication requires pair-wise action – one sender, one receiver
- No syntax control
- All data has to be exchanged – explicitly
- SPMD (single program multiple data) style
- Different work to different tasks organized by their unique ID
- No automatic work sharing
- Tasks work completely independent of each other
- Explicit synchronization required

# Point-to-Point

```
MPI_Send(&what, length, type, adressat, tag, comm);  
MPI_Recv(&what, length, type, sender, tag, comm, &status);
```

2 types of information are required to send/receive data:

- data specific (memory location, length, type)
- package information (adress, identification possibilities)



## envelope

- stamp = tag
- form and shape = communicator

# Point-to-Point

## data specific information:

- Address in memory (&what)
- Length (counted in words – not bytes)
- Type
  - pre-defined like MPI\_Int, MPI\_Double
  - user-defined (not covered in this lecture)
- neither length nor type do have to match
- receive buffer has to be large enough to keep all data sent

send 2 doubles (2x64)  
receive 4 float (4x32)

ok

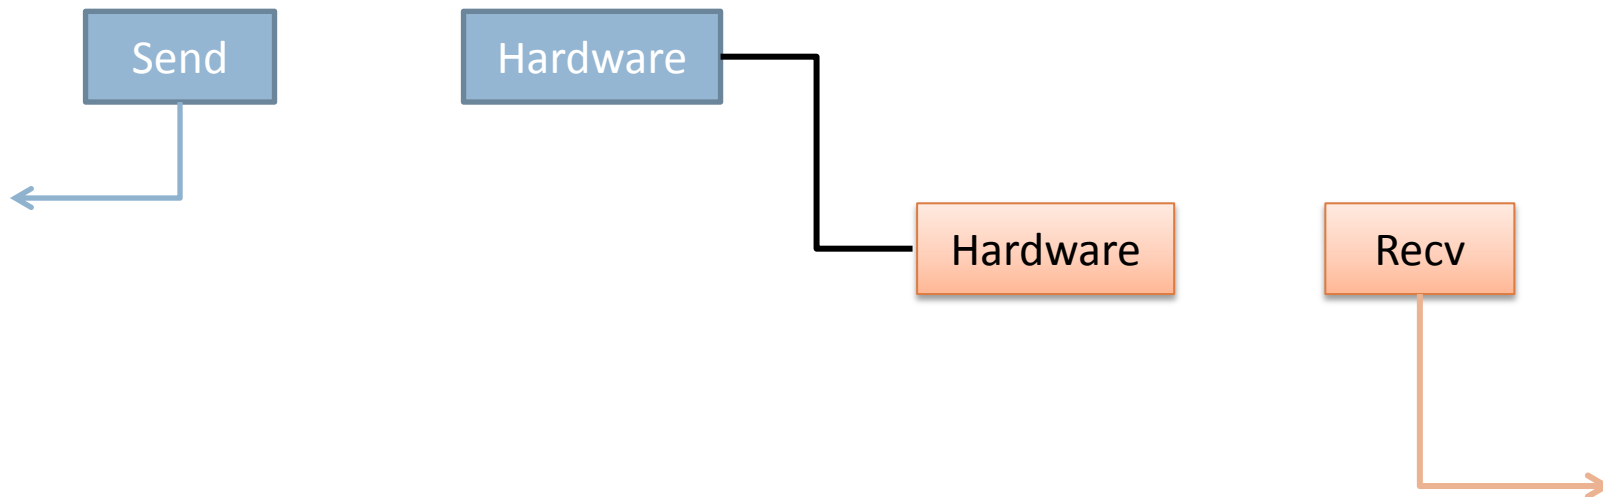
send 2 doubles (2x64)  
receive 2 integers (2x32)

not ok

# Send / Recv

What actually happens at MPI\_Send and MPI\_Recv?

Players:



# blocking communication

Actual implementation of MPI\_Send is undefined, except that it is **blocking**.

## **Blocking communication:**

memory location employed in communication is after return of call  
**(re)usable**


```
x=1.;  
MPI_Send (&x,1,MPI_DOUBLE,...);  
x=4.;  
MPI_Receive(&y,1,MPI_DOUBLE, ....);  
z=y;
```

will give correct results. x=1 is sent and x=y.



# Send / Recv

What actually happens at MPI\_Send and MPI\_Recv?



# non-blocking communication

```
x=1.;  
MPI_Isend (&x,1,MPI_DOUBLE,...);  
x=4.;
```


A non-blocking MPI\_Isend returns immediately. So the value sent may either be 1. or 4.!

```
MPI_Irecv (&x,1,MPI_DOUBLE,...);  
y=x;
```

A nonblocking MPI\_Irecv returns immediately. So the value of y may equal x or not.

# Isend / Irecv

What actually happens at MPI\_Send nad MPI\_Recv?



# blocking communication

Synchronous

**MPI\_Ssend(...)**

At return, the receiver has (at least) started to receive the message

Buffered

**MPI\_Bsend(...)**

Message is copied before sending. At return, copying is finished.

Rendezvous

**MPI\_Rsend(...)**

Receive has to be issued in advance. At return, message is as good as delivered

Why that many different Sends?

# blocking communication

## Situation:

```
if (myrank%2==0)    neighbor=myrank+1;  
else                neighbor=myrank-1;
```

```
MPI_Send(..., neighbor, ...);  
MPI_Receive(..., neighbor, ...)
```



## Deadlock:

Mutual waiting – forever (until time limit is reached)

```
if (myrank%2==0)    neighbor=myrank+1;  
else                neighbor=myrank-1;
```

```
MPI_Bsend(..., neighbor, ...);  
MPI_Receive(..., neighbor, ...)
```

Both copy and  
proceed to receive

# blocking communication

## Situation:

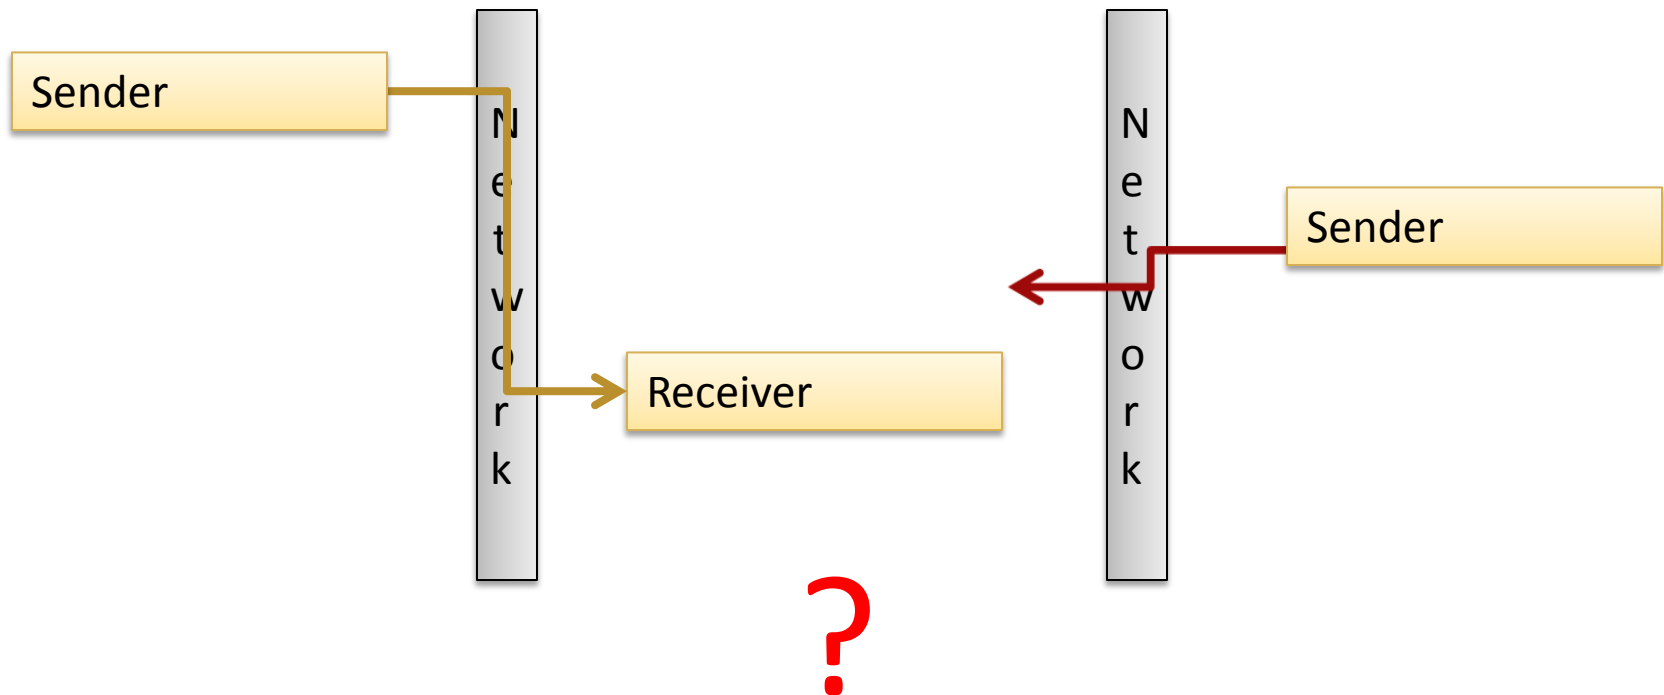
```
N=100 000 000;  
MPI_Send(Buffer,N,MPI_INT,neighbor, ...);
```



Everything gets stuck, if the huge message is not deliverable.  
MPI\_Rsend guarantees that receiver is ready.



# Using wildcards



Be careful when using MPI\_ANY\_TAG, MPI\_ANY\_SOURCE

# nonblocking communication

all blocking communications exist in a nonblocking version

MPI\_Isend  
MPI\_Ibsend  
MPI\_Irsend  
MPI\_Issend

I is synonymous for immediate

What are nonblocking send operations good for?

# Example


Write a client-server application.

**Problem:** Find best fitting parameters out of a set of 50 parameters.


**Solution** (genetic algorithm):

1. Try 1000 independent combinations of parameters (called a generation).
2. Select the 10 best fitting combinations (called evolutionary selection).
3. Generate 1000 independent combinations based on the 10 best fitting (called mutation).
4. Go back to step 1 until result is not changing any more.

generation 1



generation 2





# Example

server – clients type of code.

One master does the bookkeeping, collects results and organizes generations.  
Slaves do work and report to master.

Which slave will finish it's work in which order?

How to prevent, that messages from hundreds of slaves block the network card?

One answer – nonblocking receive on the masters side.

# nonblocking receive

```
...
MPI_Comm_rank(comm, &myrank);                                /* get ID */
MPI_Comm_size(comm, &anz);                                    /* number of Tasks */
if(myrank==0) {                                              /* specify master task */
    /* use as many receive buffers as there are tasks */
    for(i=1; i<anz; i++)
        MPI_Irecv(&puffer[i], MAX, MPI_DOUBLE, i, tag, comm, &req[i]);
    i=1;
    for(j=1; j<anz; j++) {
        for(info=0; !info; i=i%(anz-1)+1)                    /* data received? */
            MPI_Test(&req[i], &info, &status);
        MPI_Get_count(&status, MPI_DOUBLE, &ndat);           /* how much */
        do_work(&puffer[i], ndat);                            /* evaluate, set up new parameters */
    }
} else {                                                      /* slave tasks */
    evaluate_parameter_combination(quality, &ndat);
    MPI_Send(quality, ndat, MPI_DOUBLE, 0, tag, comm);
}
...
```

# usage of nonblocking sending

```
MPI_Isend(...&request);  
set_timer(&start);  
.  
get_timer(&jetzt);  
MPI_Test(&request, &flag, &status);  
if(!flag && jetzt-start>erlaubt) {  
    printf(„Receiver is dead - use another one\n");  
    .  
}
```

- blocking send operations get stuck if messages are not received – whole application has to be killed.
- nonblocking send operations allow the introduction of a timer and appropriate actions



# non blocking completion

Non blocking communication requires synchronization.

```
MPI_Wait(&request, &status)
```

(blocking) wait until communication is finished

```
MPI_Test(&request, &flag, &status)
```

look whether communication has finished. Flag is true (1) if finished, else false (0).



Einführung in das Hochleistungsrechnen  
Introduction to High Performance Computing

**VIELEN DANK**  
**THANK YOU**