



# Introduction to High Performance Computing

## Performance of OpenMP




# Performance

## Microbenchmark - Bull, Edinburgh

<https://www.epcc.ed.ac.uk/research/computing/performance-characterisation-and-benchmarking/epcc-openmp-micro-benchmark-suite>

1. Measures overhead of parallel directives and primitives
2. Measures overhead of different scheduling strategies

# Performance




# Performance

Tests were executed on 4 cores with hyperthreading (8 virtual cores)

## Results:

- No extra penalty for virtual cores (except for barrier and for)
- The more threads, the more overhead
- Reduction is very expensive
- PARALLEL is cheaper than PARALLEL FOR, but this is cheaper than a PARALLEL and a FOR. **One loop, use PARALLEL FOR**
- **PARALLEL is expensive** – use a parallel region as long as possible
- FOR is as cheap as the barrier at its end – use NOWAIT
- SINGLE and CRITICAL are very cheap synchronizations and show only little dependency on the number of threads used

# Performance





# Performance

Test with 4 Threads on a 2x4 core machine.

**Blue:** Threads are located on one socket (occupy one CPU)


**Yellow:** Threads are spread – two on each socket

Collocating threads is very important (LSF does this automatically if possible)

Possibilities to bind threads to cores:

- KMP\_AFFINITY ([Intel](#))
- Environment variable OMP\_PROC\_BIND={true, false} (a thread should not be moved by OS)
- GOMP\_CPU\_AFFINITY in conjunction with GNU-compilers

# Performance







# Performance

Different scheduling strategies

## **Results:**

- STATIC shows least overhead irrespective of chunk size
- DYNAMIC with very small chunks (1) should be avoided if possible
- GUIDED generates considerable overhead – only to be used if really necessary





# False Sharing

Considering multi-core applications caches and cache coherency gain increasing importance.

The following examples shows 2 version of an alternative to a reduction. The first one with automatic variables allocated on the stack of each thread, the second one with a shared array there each thread access the corresponding array element.

# Examples

```
#define Intervalle 10000
void main(int argc, char **argv) {
    int i;
    double step,x,pi,sum=0.,psum;
    step=1./Intervalle;
    /* Fork and work sharing combined */
    #pragma omp parallel shared(sum)
    {
        #pragma omp for private(x,i,psum) nowait
        for(i=0,psum=0.;i<Intervalle;i++) {
            x=(i+0.5)*step;
            psum+=4.0/(1.0+x*x);
        }
        #pragma omp critical
        sum+=psum;
    } /* end of parallel region */
    pi=step*sum;
}
```

# Examples

```
#define Intervalle 10000
void main(int argc, char **argv) {
    int i;
    double step,x,pi,sum=0.,psum[max_threads];
    step=1./Intervalle;
    /* Fork and work sharing combined */
    #pragma omp parallel shared(sum)
    {
        me=omp_get_thread_num();
        #pragma omp for private(x,i)
        for(i=0,psum[me]=0.;i<Intervalle;i++) {
            x=(i+0.5)*step;
            psum[me]+=4.0/(1.0+x*x);
        }
        #pragma omp single
            for(i=0;i<omp_get_team_size();i++) sum+=psum[i];
    } /* end of parallel region */
    pi=step*sum;
}
```



# False Sharing

- Caches are organized in cache lines – usually one line 2 words.
- `psum[4]` may occupy two cache lines when loaded to a core.
- Is `psum[1]` changed, the cache coherency protocol of the system invalidates the cache line in all other caches.
- Thread 0 has to reload `psum[0]` (and the complete cache line).
- Is one of the values changed, the cache line is invalidated .....

**Don't use small arrays for writing by different threads.**



Einführung in das Hochleistungsrechnen  
Introduction to High Performance Computing

**VIELEN DANK**  
**THANK YOU**